

```
In [1]: # @hidden_cell
# The project token is an authorization token that is used to access
# project resources like data sources, connections, and used by platfo
# rm APIs.
from project_lib import Project
project = Project(project_id='ca8998af-60ec-4af7-a5ee-935daeff58ca',
project_access_token='p-1379f525698dc66ee90bcd6b8164145eb35ea2e7')
pc = project.project_context
```

```
In [2]: # @hidden_cell
! pip install arch
```

```
Requirement already satisfied: arch in /opt/conda/envs/Python-3.7-
main/lib/python3.7/site-packages (4.15)
Requirement already satisfied: property-cached>=1.6.3 in /opt/cond
a/envs/Python-3.7-main/lib/python3.7/site-packages (from arch) (1.
6.4)
Requirement already satisfied: statsmodels>=0.9 in /opt/conda/envs
/Python-3.7-main/lib/python3.7/site-packages (from arch) (0.11.1)
Requirement already satisfied: scipy>=1.0.1 in /opt/conda/envs/Pyt
hon-3.7-main/lib/python3.7/site-packages (from arch) (1.5.0)
Requirement already satisfied: cython>=0.29.14 in /opt/conda/envs/
Python-3.7-main/lib/python3.7/site-packages (from arch) (0.29.21)
Requirement already satisfied: pandas>=0.23 in /opt/conda/envs/Pyt
hon-3.7-main/lib/python3.7/site-packages (from arch) (1.0.5)
Requirement already satisfied: numpy>=1.14 in /opt/conda/envs/Pyth
on-3.7-main/lib/python3.7/site-packages (from arch) (1.18.5)
Requirement already satisfied: patsy>=0.5 in /opt/conda/envs/Pytho
n-3.7-main/lib/python3.7/site-packages (from statsmodels>=0.9->arc
h) (0.5.1)
Requirement already satisfied: pytz>=2017.2 in /opt/conda/envs/Pyt
hon-3.7-main/lib/python3.7/site-packages (from pandas>=0.23->arch)
(2020.1)
Requirement already satisfied: python-dateutil>=2.6.1 in /opt/cond
a/envs/Python-3.7-main/lib/python3.7/site-packages (from pandas>=
0.23->arch) (2.8.1)
Requirement already satisfied: six in /opt/conda/envs/Python-3.7-m
ain/lib/python3.7/site-packages (from patsy>=0.5->statsmodels>=0.9
->arch) (1.15.0)
```

Importation des Bilbiothèques

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
sns.set()
```

Defintion des fonctions des graphes

```
In [116]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller
def tsplot(y, ADF=True, lags=None, title=None, figsize=(15, 7), style
='bmh'):
    """
    Plot time series, its ACF and PACF, calculate Dickey-Fuller
    test

    y - timeseries
    lags - how many lags to include in ACF, PACF calculation
    """
    if title == None:
        title = "Time Series Analysis Plots"
    else:
        title = str(title)
    if not isinstance(y, pd.Series):
        y = pd.Series(y)

    with plt.style.context(style):
        fig = plt.figure(figsize=figsize)
        layout = (2, 2)
        ts_ax = plt.subplot2grid(layout, (0, 0), colspan=2)
        acf_ax = plt.subplot2grid(layout, (1, 0))
        pacf_ax = plt.subplot2grid(layout, (1, 1))

        y.plot(ax=ts_ax)
        if ADF == True:
            p_value = adfuller(y)[1]
            ts_ax.set_title(title + '\n Dickey-Fuller: p_value =
{0:.5f}'.format(p_value))
        else:
            ts_ax.set_title(title)
        plot_acf(y, lags=lags, ax=acf_ax)
        plot_pacf(y, lags=lags, ax=pacf_ax)
        plt.tight_layout()
```

```
In [128]: from statsmodels.graphics.gofplots import qqplot
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller
from seaborn import distplot
def tsplot_resid(y, lags=None, title=None, figsize=(15, 7), style='bmh'):
    """
    Plot time series, its ACF and PACF, calculate Dickey-Fuller test

    y - timeseries
    lags - how many lags to include in ACF, PACF calculation
    """
    if title == None:
        title = "Time Series Analysis Plots"
    else:
        title = str(title)
    if not isinstance(y, pd.Series):
        y = pd.Series(y)

    with plt.style.context(style):
        fig = plt.figure(figsize=figsize)
        layout = (3, 2)
        ts_ax = plt.subplot2grid(layout, (0, 0), colspan=2)
        acf_ax = plt.subplot2grid(layout, (1, 0))
        pacf_ax = plt.subplot2grid(layout, (1, 1))
        qqplot_ax = plt.subplot2grid(layout, (2, 0))
        dist_ax = plt.subplot2grid(layout, (2, 1))

        y.plot(ax=ts_ax)
        ts_ax.set_title(title)
        plot_acf(y, lags=lags, ax=acf_ax)
        plot_pacf(y, lags=lags, ax=pacf_ax)
        qqplot(y, fit=True, line="45", ax=qqplot_ax)
        qqplot_ax.set_title("Normalité")
        distplot(y, ax=dist_ax)
        dist_ax.set_title("Distribution")
        plt.tight_layout()
```

Importation des données

```
In [75]: consommation = project.get_file("Consommation d'électricité.xls")
df = pd.read_excel(consommation, sheet_name = "Algerie", index_col=
0)
df = df.transpose()
df.head()
```

Out[75]:

Année	Consommation d'electricité par habitant
1971	133.873489
1972	142.875928
1973	158.754155
1974	170.660458
1975	195.692277

Traitement des données

```
In [35]: df["Année"] = df.index
df.reset_index(inplace=True)
df["Année"] = pd.to_datetime(df["Année"], format="%Y")
df.drop(["index"], axis=1, inplace=True)
df.set_index("Année", inplace = True)
df.head()
```


Out[35]:

Année	Consommation d'electricité par habitant
Année	
1971-01-01	133.873489
1972-01-01	142.875928
1973-01-01	158.754155
1974-01-01	170.660458
1975-01-01	195.692277

Graphique de la serie chronologique


```
In [109]: df.plot(figsize=(15,7))
plt.xlabel("Années")
plt.ylabel("Consommation d'electricité par habitant")
plt.title("Consommation d'electricité par habitant par année")
```

```
Out[109]: Text(0.5, 1.0, "Consommation d'electricité par habitant par année")
```



Examen des données

```
In [119]: tsplot(df["Consommation d'electricité par habitant"],lags=25, title
= "Serie avant integration (differentiation)")
```



- Après avoir appliqué le test de Dickey-Fuller Augmenté, la $p_value = 0.99898 > 0.05$, ce qui signifie que les données ne sont pas stationnaires.


Differentiation des données

```
In [13]: dff = df.diff()
dff = dff.dropna()
dff.head()
```

Out[13]:

Année	Consommation d'électricité par habitant
Année	
1972-01-01	9.002438
1973-01-01	15.878228
1974-01-01	11.906302
1975-01-01	25.031819
1976-01-01	24.373496

```
In [120]: tsplot(dff["Consommation d'électricité par habitant"],lags=25, title
="Serie après integration d'ordre 1")
```



- La nouvelle série après différenciation est une série stationnaire selon la p_value du test ADF qui est égale à $0.007 < 0.05$
- Aucun pic n'est significatif ni dans la FAC ni dans la FAP
- À première vue, un processus ARIMA(0,1,0) est le meilleur pour cette série de données

```
In [16]: from statsmodels.tsa.arima_model import ARIMA
def _get_best_model(TS):
    best_aic = np.inf
    best_order = None
    best_mdl = None
    pq_rng = range(5) # [0,1,2,3,4]
    d_rng = range(2) # [0,1]
    for i in pq_rng:
        for d in d_rng:
            for j in pq_rng:
                try:
                    tmp_mdl = ARIMA(TS, order=(i,d,j)).fit()
                    tmp_aic = tmp_mdl.aic
                    if tmp_aic < best_aic:
                        best_aic = tmp_aic
                        best_order = (i, d, j)
                        best_mdl = tmp_mdl
                except: continue
    print('aic: {:.6.2f} | order: {}'.format(best_aic, best_order))
    return best_aic, best_order, best_mdl
res_tup = _get_best_model(df)
```

```
aic: 488.37 | order: (0, 1, 0)
```

Choix du meilleure modèle

```
In [123]: def get_best_model_2(TS):
    AIC = []
    order = []
    pq_rng = range(3) # [0,1,2,3,4]
    d_rng = range(1,2) # [0,1]
    for d in d_rng:
        for i in pq_rng:
            for j in pq_rng:
                try:
                    tmp_mdl = ARIMA(TS, order=(i,d,j)).fit()
                    AIC.append(tmp_mdl.aic)
                    order.append((i,d,j))
                except: continue
    tbl = {"Ordre":order, "AIC":AIC}
    rank = pd.DataFrame(tbl)
    rank.sort_values(by=["AIC"], ascending=True, inplace=True)
    rank.set_index("Ordre", inplace=True)
    return rank
```

```
In [124]: get_best_model_2(df)
```

```
Out[124]:
```

AIC	
Ordre	
(0, 1, 0)	488.371799
(2, 1, 0)	489.413460
(1, 1, 2)	489.867997
(0, 1, 2)	489.879725
(1, 1, 0)	489.972328
(2, 1, 2)	490.062727
(0, 1, 1)	490.093880
(2, 1, 1)	490.167045
(1, 1, 1)	492.251340

- Selon le critère Akaike (AIC), le meilleur modèle est un ARIMA(0,1,0) avec un AIC = 488.37
- Cela confirme notre première hypothèse

Estimation du modèle ARIMA(0,1,0)

```
In [125]: modele = ARIMA(df, order=(0,1,0)).fit()
modele.summary()
```

```
Out[125]:
```

ARIMA Model Results

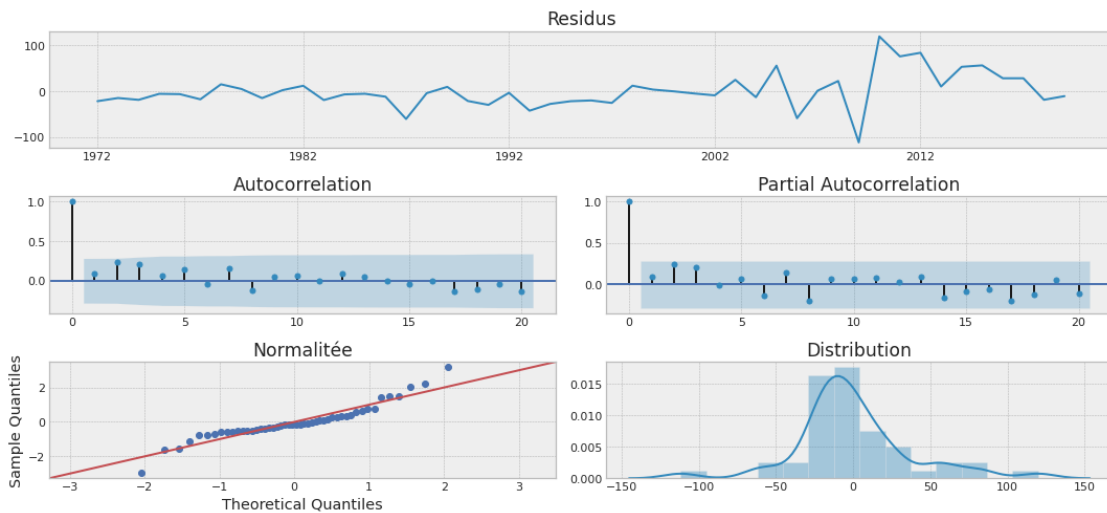
Dep. Variable:	D.Consommation d'électricité par habitant	No. Observations:	48
Model:	ARIMA(0, 1, 0)	Log Likelihood	-242.186
Method:	css	S.D. of innovations	37.585
Date:	Mon, 28 Dec 2020	AIC	488.372
Time:	22:08:33	BIC	492.114
Sample:	01-01-1972	HQIC	489.786
	- 01-01-2019		

	coef	std err	z	P> z	[0.025	0.975]
const	30.5443	5.425	5.630	0.000	19.912	41.177

- La constante du modèle est significative

Examen des Résidus


```
In [129]: tsplot_resid(modele.resid, lags=20, title="Residus")
```



- A première vue, les résidus ont les propriétés d'un Bruit Blanc
- Il n'existe aucun pic significatif dans la FAC et la FAP

Test ARCH d'Engle pour l'homoscédasticité des Erreurs

H₀ : Les résidus au carré sont une séquence de bruit blanc - les résidus sont homoscédastiques.

H₁ : Les résidus au carré n'ont pas pu être ajustés avec un modèle de régression linéaire et présentent une hétéroscédasticité.

```
In [130]: from statsmodels.stats.diagnostic import het_arch
print()
print("La valeur critique :", het_arch(modele.resid) [2], ",          la p_
value :", het_arch(modele.resid) [3])
```

```
La valeur critique : 1.6320909879309025 ,          la p_value : 0.1505
1916477323907
```

- Étant donné que la p_value est de 0.1505 > 0.05, nous ne parvenons pas à rejeter l'hypothèse nulle, d'où les résidus sont homoscédastiques
- Donc il n'existe pas d'effet ARCH sur les résidus.

Test de Ljung-Box pour Autocorrelation des Erreurs

H₀: Les résidus sont distribués indépendamment.

H_A: les résidus ne sont pas distribués indépendamment; ils présentent une corrélation en série.

```
In [134]: from statsmodels.stats.diagnostic import acorr_ljungbox
LB = acorr_ljungbox(modele.resid, lags=[10])
print("statistique =", LB[0][0], "                                p_value =", LB[1][0])
)
```

statistique = 10.181333984986784 p_value = 0.424730
9703532606

- Étant donné que la p_value est de $0,4247 > 0,05$, nous ne parvenons pas à rejeter l'hypothèse nulle, les résidus sont donc distribués indépendamment,
- Ce qui signifie que les résidus ne sont pas autocorrélés

Test de Jarque-Bera de Normalité des Erreurs

H0 : les données suivent une loi normale.

H1 : les données ne suivent pas une loi normale.

```
In [137]: from scipy.stats import jarque_bera
JB = jarque_bera(modele_garch.resid)
print("statistic=", JB[0], "                                p_value=", JB[1])
print()
```

statistic= 12.584370443790368 p_value= 0.00185071131000
53596

- Puisque la $p_value = 0.001 < 0,05$, nous rejetons l'hypothèse nulle. Ainsi, nous avons des preuves suffisantes pour dire que ces données ont une asymétrie et un kurtosis qui sont significativement différents d'une distribution normale.
- Donc les erreurs ne suivent pas une loi Normale


Prevision pur les années 2020 - 2024

```
In [138]: pred = modele.forecast(5)
années = [2020, 2021, 2022, 2023, 2024]
df_pred = pd.DataFrame({"Années":années, "Prevision":pred[0]})
df_pred.set_index("Années", inplace=True)
df_pred
```

Out[138]:

Prevision	
Années	
2020	1630.544302
2021	1661.088605
2022	1691.632907
2023	1722.177209
2024	1752.721512

```
In [139]: fig, ax = plt.subplots(figsize=(15, 7))
fig = modele.plot_predict(1,55, ax=ax)
```



ARCH / GARCH


Fractionnement des données en données d'entraînement et de test

```
In [7]: size = int(len(df)*0.9)
df_train = df.iloc[:size]
df_test = df.iloc[size:]
df_train.tail()
```

Out[7]:

Année	Consommation d'électricité par habitant
Année	
2010-01-01	1016.636669
2011-01-01	1123.332731
2012-01-01	1237.966507
2013-01-01	1278.915343
2014-01-01	1362.871919

```
In [113]: tsplot(df_train["Consommation d'électricité par habitant"],lags=25,
title = "Serie avant integration (differentiation)")
```




```
In [12]: dif_train = df_train.diff()
dif_train = dif_train.dropna()
dif_train.head()
```

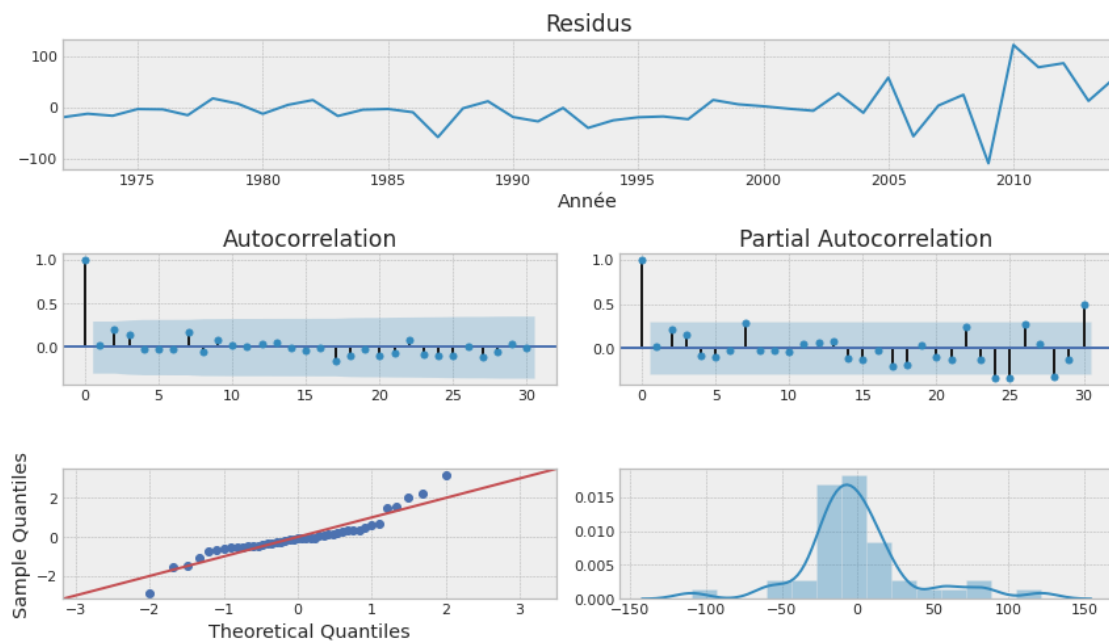
Out[12]:

Année	Consommation d'électricité par habitant
Année	
1972-01-01	9.002438
1973-01-01	15.878228
1974-01-01	11.906302
1975-01-01	25.031819
1976-01-01	24.373496

```
In [14]: tsplot(dif_train["Consommation d'électricité par habitant"],lags=30,
title="Serie après integration d'ordre 1")
```



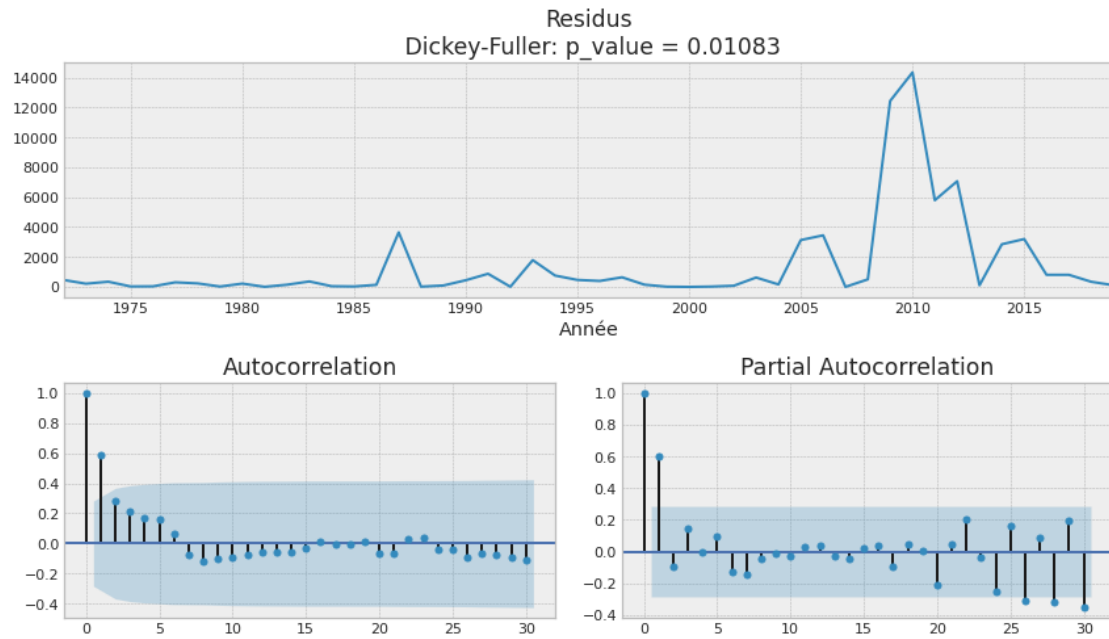
```
In [45]: model = ARIMA(df_train, order=(0,1,0)).fit()
tsplot_resid(model.resid,lags=30, title="Residus")
```



```
In [ ]:
```

Elevation des residus au carré pour detecter la necessité d'un modèle ARCH pour les residus

```
In [50]: residus_carré = modele.resid**2  
         tsplot(residus_carré, lags=30, title="Residus")
```



- Le carré des résidus n'est pas un Bruit Blanc à cause de la présence de pics significatifs dans la FAC et la FAP
- Selon les deux graphes des autocorrelation, un modèle GARCH(1,1) est le meilleur modèle pour expliquer la volatilité dans les résidus

Estimation d'un modèle GARCH

```
In [51]: from arch import arch_model
garch = arch_model(modele.resid,p=1, q=1)
modele_garch = garch.fit(update_freq=5)
modele_garch.summary()
```

```
Iteration:      5,      Func. Count:      28,      Neg. LLF: 231.17058190
25867
Iteration:     10,      Func. Count:      53,      Neg. LLF: 231.08722051
3383
Iteration:     15,      Func. Count:      78,      Neg. LLF: 231.07867542
736687
Iteration:     20,      Func. Count:     103,      Neg. LLF: 230.99550939
67152
Optimization terminated successfully      (Exit mode 0)
      Current function value: 230.9949613661152
      Iterations: 24
      Function evaluations: 122
      Gradient evaluations: 24
```

Out[51]:

Constant Mean - GARCH Model Results

Dep. Variable:	None	R-squared:	-0.060
Mean Model:	Constant Mean	Adj. R-squared:	-0.060
Vol Model:	GARCH	Log-Likelihood:	-230.995
Distribution:	Normal	AIC:	469.990
Method:	Maximum Likelihood	BIC:	477.475
		No. Observations:	48
Date:	Mon, Dec 28 2020	Df Residuals:	44
Time:	20:11:23	Df Model:	4

Mean Model

	coef	std err	t	P> t	95.0% Conf. Int.
mu	-9.2138	4.516	-2.040	4.132e-02	[-18.065, -0.363]

Volatility Model

	coef	std err	t	P> t	95.0% Conf. Int.
omega	189.8474	117.529	1.615	0.106	[-40.505, 4.202e+02]
alpha[1]	0.3799	0.156	2.433	1.497e-02	[7.385e-02, 0.686]
beta[1]	0.4789	0.185	2.593	9.503e-03	[0.117, 0.841]

Covariance estimator: robust

```
In [53]: df_train2 = df_train**2
df_train2.head()
```

Out[53]:


Année	Consommation d'électricité par habitant
Année	
1971-01-01	17922.111115
1972-01-01	20413.530678
1973-01-01	25202.881870
1974-01-01	29124.991867
1975-01-01	38295.467311

```
In [52]: df2 = df**2
df2.head()
```


Out[52]:

Année	Consommation d'électricité par habitant
Année	
1971-01-01	17922.111115
1972-01-01	20413.530678
1973-01-01	25202.881870
1974-01-01	29124.991867
1975-01-01	38295.467311


```
In [54]: tsplot(df2["Consommation d'électricité par habitant"],lags=19, title
="Serie après integration d'ordre 1")
```





```
In [55]: tsplot_resid(modele.resid)
```



```
In [56]: tsplot_resid(modele_garch.resid)
```



```
In [57]: garch2 = modele_garch.resid**2
         tsplot_resid(garch2)
```



```
In [135]: AIC = []
          order = []
          for p in range(1,11):
              for q in range(2):
                  garch = arch_model(model.resid,p=p, q=q)
                  modele_garch = garch.fit(update_freq=50)
                  AIC.append(tmp_mdl.aic)
                  order.append((p,q))

          tbl = {"Ordre":order, "AIC":AIC}
          df = pd.DataFrame(tbl)
          df
```

```
Optimization terminated successfully      (Exit mode 0)
      Current function value: 229.40434518955567
      Iterations: 16
      Function evaluations: 67
      Gradient evaluations: 16
```

```
-----
NameError                                Traceback (most recent c
all last)
<ipython-input-135-c903177ca88b> in <module>
      5         garch = arch_model(model.resid,p=p, q=q)
      6         modele_garch = garch.fit(update_freq=50)
----> 7         AIC.append(tmp_mdl.aic)
      8         order.append((p,q))
      9

NameError: name 'tmp_mdl' is not defined
```

```
In [136]: df.sort_values(by="AIC", ascending=True, inplace=True)
df.head(1)
```

Out[136]:

```

      AIC
Ordre
(0, 1, 1) 479.129958
```

```
In [149]: garch = arch_model(model.resid, vol='GARCH', p=1, q=1)
modele_garch = garch.fit(update_freq=50)
modele_garch.summary()
```

```

Optimization terminated successfully      (Exit mode 0)
      Current function value: 227.68691210273192
      Iterations: 23
      Function evaluations: 117
      Gradient evaluations: 23
```

Out[149]:

Constant Mean - GARCH Model Results

```

Dep. Variable:          None      R-squared:    -0.013
Mean Model:      Constant Mean  Adj. R-squared: -0.013
Vol Model:       GARCH          Log-Likelihood: -227.687
Distribution:    Normal          AIC:         463.374
Method: Maximum Likelihood      BIC:         470.774

                                No. Observations:    47
Date:      Mon, Dec 28 2020      Df Residuals:    43
Time:      18:32:33              Df Model:       4
```

Mean Model

```

      coef  std err      t  P>|t|  95.0% Conf. Int.
mu  -6.0977   4.910  -1.242  0.214  [-15.721, 3.526]
```

Volatility Model

```

      coef  std err      t      P>|t|    95.0% Conf. Int.
omega  176.5956  114.559   1.542    0.123  [-47.935,4.011e+02]
alpha[1]  0.2931   0.136   2.153  3.132e-02  [2.628e-02, 0.560]
beta[1]  0.5872   0.129   4.557  5.179e-06  [ 0.335, 0.840]
```

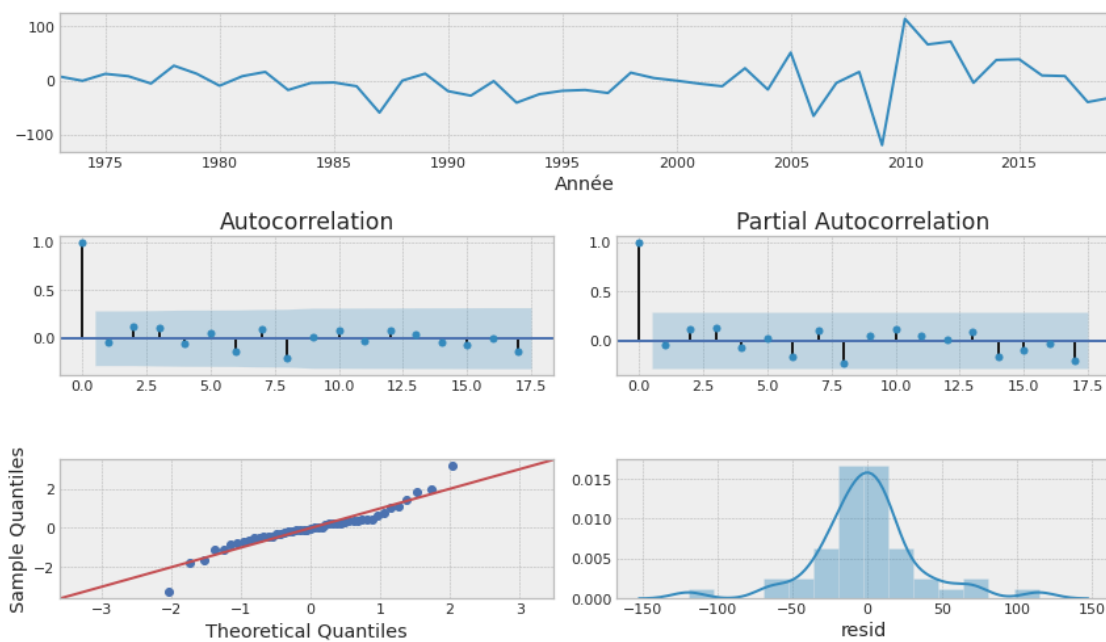
Covariance estimator: robust

```
In [138]: print("La valeur critique :",het_arch(modele_garch.resid)[0],het_arch(modele_garch.resid)[2] ,",
la p_value :",het_arch(modele_garch.resid)[1],het_arch(modele_garch.resid)[3])
```

```

La valeur critique : 14.583755980836452 1.6915307273492852 ,
la p_value : 0.1479902669312523 0.13637259886661823
```

```
In [143]: tsplot_resid(modele_garch.resid)
```



```
In [59]: AIC = []
order = []
p_val = []
for p in range(1,11):
    for q in range(2):
        garch = arch_model(modele.resid, vol='GARCH', p=p, q=q)
        modele_garch = garch.fit(update_freq=50)
        AIC.append(modele_garch.aic)
        order.append((p,q))
        p_val.append(het_arch(modele_garch.resid)[3])
tbl = {"Ordre":order, "P_val":p_val, "AIC":AIC}
df = pd.DataFrame(tbl)
df.sort_values(by="P_val", ascending=True, inplace=True)
```

```
Optimization terminated successfully      (Exit mode 0)
Current function value: 232.9765195734206
Iterations: 24
Function evaluations: 97
Gradient evaluations: 24
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.9949613661152
Iterations: 24
Function evaluations: 122
Gradient evaluations: 24
Optimization terminated successfully      (Exit mode 0)
Current function value: 231.46459894576708
Iterations: 26
Function evaluations: 132
Gradient evaluations: 26
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.99496121007604
Iterations: 27
Function evaluations: 166
Gradient evaluations: 27
Optimization terminated successfully      (Exit mode 0)
Current function value: 231.12671119026473
Iterations: 29
Function evaluations: 178
Gradient evaluations: 29
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.9464724923102
Iterations: 28
Function evaluations: 203
Gradient evaluations: 28
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531736435568
Iterations: 34
Function evaluations: 245
Gradient evaluations: 34
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.0053176639155
Iterations: 35
Function evaluations: 288
Gradient evaluations: 35
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.0053174436408
Iterations: 31
Function evaluations: 251
Gradient evaluations: 31
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531733630223
Iterations: 31
Function evaluations: 283
Gradient evaluations: 31
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531732748857
Iterations: 36
Function evaluations: 330
Gradient evaluations: 36
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531735829298
Iterations: 29
Function evaluations: 293
Gradient evaluations: 29
```

```
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.0053174873644
Iterations: 34
Function evaluations: 345
Gradient evaluations: 34
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531735025632
Iterations: 30
Function evaluations: 335
Gradient evaluations: 30
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531768885725
Iterations: 35
Function evaluations: 390
Gradient evaluations: 35
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531760298708
Iterations: 30
Function evaluations: 365
Gradient evaluations: 30
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531740228558
Iterations: 33
Function evaluations: 401
Gradient evaluations: 33
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531780911427
Iterations: 30
Function evaluations: 395
Gradient evaluations: 30
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531735473038
Iterations: 31
Function evaluations: 409
Gradient evaluations: 31
Optimization terminated successfully      (Exit mode 0)
Current function value: 230.00531741232135
Iterations: 28
Function evaluations: 397
Gradient evaluations: 28
```

```
In [146]: df.sort_values(by="AIC", ascending=True, inplace=True)
df
```

Out[146]:

	Ordre	P_val	AIC
4	(3, 0)	0.192056	460.602458
6	(4, 0)	0.192272	462.601188
5	(3, 1)	0.192069	462.602458
1	(1, 1)	0.136373	463.373824
7	(4, 1)	0.192263	464.601188
8	(5, 0)	0.192271	464.601188
0	(1, 0)	0.114131	464.808690
3	(2, 1)	0.136374	465.373825
10	(6, 0)	0.192248	466.601188
9	(5, 1)	0.191536	466.602406
2	(2, 0)	0.100641	466.703509
12	(7, 0)	0.192249	468.601188
11	(6, 1)	0.191530	468.602407
14	(8, 0)	0.192258	470.601188
13	(7, 1)	0.191534	470.602407
16	(9, 0)	0.197059	472.482455
15	(8, 1)	0.191533	472.602408
17	(9, 1)	0.197059	474.482455
18	(10, 0)	0.197064	474.482455
19	(10, 1)	0.197039	476.482455

In []: